

On Extensibility and Personalizability of Multi-Modal Trip Planning

Abstract

Trip planning is a practical task that has drawn extensive attention from the AI planning and scheduling community and the industrial/commercial sectors. In this paper, we consider the setting of the *multi-modal* trip planning, where users can exploit different transportation modes, such as walking, biking, public transit, and taxi. In such a context, it would be of benefit if the user was able to extend the *cost base*, including traveling time and fare, of the planner, and to personalize the planner according to her own constraints and preferences. To this end, we designed and developed a hypergraph-based multi-modal trip planner that allows users to upload auxiliary cost metrics (e.g., crime rates), to specify constraints as a theory in the *linear temporal logic*, and to express preferences as a *preferential cost function*.

Introduction

Trip planning, as an application of planning and scheduling, has seen substantive implementations by researchers and developers (Bast et al. 2015). Some of the planning systems are *multi-modal*; that is, combining distinct transportation modes, the trip planners compute optimal routes from sources to destinations. This notion of “optimality” generally refers to the computed routes having minimal total time or total fare.

However, in the eyes of a user, it may be more faceted than just “fastest” or “cheapest.” For instance, for a college student who specifies that she will only walk or take public transit in a trip from Palo Alto to San Francisco, the computed plan is not necessarily the fastest (e.g., taking a cab could be faster) or the cheapest (e.g., walking all the way has no fare). This happens when a user tells the planner her hard constraints, called *constraints*. The planner then needs to either satisfy the constraints in the search process or return failure because they are over-restrictive. Moreover, the user might want to further customize the planner by describing soft constraints, called *preferences*. For example, an agent wants to travel from school to downtown, and prefers biking to taking the bus. Thus, a trip with more biking than bus may be considered better for the agent than the one with more bus than biking. In this case, the planner will need to accom-

modate user preferences whenever possible in the search of optimal solutions.

Representing and reasoning about constraints and preferences are fundamental to decision making in automated planning and scheduling in artificial intelligence. Son and Pontelli presented a declarative preference language, \mathcal{PP} , to represent ordinal preferences (e.g., basic desires, atomic preferences and general preferences) between trajectories of a planning problem (Son and Pontelli 2004). Bienvenu et al. proposed a first-order preference language, \mathcal{LPP} , extending \mathcal{PP} by allowing the specification of quantified qualitative preferences (Bienvenu, Fritz, and McIlraith 2011). Modesti and Sciomachen introduced an ad hoc utility function for weighing the arcs both with their time and fare (Modesti and Sciomachen 1998). This relationship between time and fare in the setting of transportation was studied, and economic models capturing the trade-off between the two was presented (Antonioni, Matsoukis, and Roussi 2007).

However, relatively limited effort has been devoted to designing and implementing real-world multi-modal trip planners that captures user constraints and preferences over the *cost base*, possibly extended from the user with auxiliary *cost metrics*, such as crime rates and pollution statistics.

Using a high-performance graph search engine (Zhou and Hansen 2011), we designed and implemented a multi-modal trip planner that uses pure graph-search. This allows us to flexibly combine various modes (e.g., walk, bike, car, public transit, and taxi) and declaratively change constraints and preferences. The planner allows the user to upload *new* mapping data over which to express constraints and preferences. For instance, a user might upload a map of crime in the city, and ask the trip planner to avoid areas where crime is frequent. Then, the planner takes constraints expressed in *linear temporal logic* to constrain the search space (e.g., never bike after transit, and never walk through bad neighborhoods). Finally, the planner uses a weighted sum, called *preferential cost function*, over several cost metrics (e.g., time spend biking, fare on public transit, and overall crimes walking through) which can be re-weighted based on different user preferences.

Our paper is organized as follows. In the next section, we present what it means for a planner to be extensible and formally define the method to incorporating new metrics into the planner. In the following section, we discuss the two

aspects of personalization in trip planning: constraints and preferences, and how they are modeled and reasoned with in the setting of multi-modal trip planning. We then move on to describe the system structure of our graph-search based planner, and show results obtained from our planner in various occasions. Finally, we conclude by outlining some future research directions.

Extensibility

Allowing users to upload their own data sets of interests is an important step towards customization of a trip planner. We designed a framework where a user can upload auxiliary cost metric data (e.g., crime statistics and pollution data) into the planner, and the planner will compute an optimal route accordingly.

The user-created data are the auxiliary data that is represented as pairs of latitude and longitude degrees. To merge these lat-long pairs into the planner, we performed a neighborhood search to calculate the total score of auxiliary data for each lat-long pair already in our planner. It might be of strong interest to some user for our planning system to take care of criminal statistics so that some level of safety of the resulting routes is guaranteed. For instance, a user traveling through the downtown area of San Francisco around midnight may want to upload a data set of crimes (cf. Figure 1), and express her constraints and preferences in hope of a safer trip plan. Note that in Figure 1 bad areas are the colored circles, whose integer labels represent numbers of crimes in corresponding areas.




Figure 1: Crime rates in San Francisco

Formally, let \mathcal{U} be the uploaded auxiliary dataset of lat-long pairs, $N = (x_N, y_N)$ a node in our planner and r the effective radius. We first compute the set $\mathcal{P} \subseteq A$ of auxiliary points such that, for every node $P \in \mathcal{P}$, the Euclidean distance between N and P is within r . Then, the score of auxiliary data for N with respect to \mathcal{P} and r , denoted by $S(N, \mathcal{P}, r)$, is computed as follows.

$$S(N, \mathcal{P}, r) = \sum_{P \in \mathcal{P}} \left(1 - \frac{\sqrt{(x_N - x_P)^2 + (y_N - y_P)^2}}{r} \right)$$

Now we turn to Figure 2 for an instance to show how auxiliary data are integrated into the graph using the equation above. In Figure 2, we have the green nodes denoting the graph nodes, and the red nodes the new auxiliary nodes that represent locations of criminal events. For node N , there are three red nodes within its neighborhood of radius of r ; thus, we have $\mathcal{P} = \{P_1, P_2, P_3\}$. We assume r is 100 feet, and the distances from N to P_1 , P_2 and P_3 are 25, 80 and 90 feet, respectively. Then, the crime score of $S(N, \mathcal{P}, r)$ is 1.05.




Figure 2: Integrating uploaded data

Personalizability

Personalizability consists of two aspects: constraints and preferences. From the viewpoint of the planner, constraints, also referred to as hard constraints, are statements that the planner has to satisfy during the planning process; whereas preferences, also called soft constraints, are specifications that the planner will need to optimize. We formulated constraints using linear temporal logic (LTL) and preferences as a preferential cost function (PCF), and implemented our planner leveraging the widely-used graph search algorithm the A*.

Constraints

As constraints in the setting of trip planning are often declarative and temporal, our choice of LTL is straightforward. We now give a brief review of linear temporal logic (LTL). Let f be a propositional formula over a finite set L of Boolean variables. LTL formulas are defined recursively as follows.

$$\varphi = f | \varphi_1 \wedge \varphi_2 | \varphi_1 \vee \varphi_2 | \neg \varphi | \bigcirc \varphi | \square \varphi | \diamond \varphi | \varphi_1 \mathcal{A} \varphi_2 \quad (1)$$

Note that we have $\varphi_1 \mathcal{A} \varphi_2$, and it means that “ φ_2 holds right after φ_1 holds.”

A natural constraint in trip planning could be “In this trip I will not drive a car after biking or taking the public transit.” In LTL, such constraint can be translated into an LTL formula

$$\psi = ((M = b) \vee (M = p)) \mathcal{A} (\square (\neg (M = c))). \quad (2)$$

As the actions in trip planning is limited to taking different transportation modes, in our definition of the se-

mantics of LTL these actions are subsumed into the interpretations of L , or *states*. The semantics of LTL is defined with regard to trajectories of states. Let σ be a trajectory of states $S_0, a_1, S_1, \dots, a_n, S_n$, and $\sigma[i]$ a suffix $S_i, a_{i+1}, S_{i+1}, \dots, a_n, S_n$. We have

$$\begin{aligned}
\sigma &\models f \text{ iff } S_0 \models f, \\
\sigma &\models \varphi_1 \wedge \varphi_2 \text{ iff } \sigma \models \varphi_1 \text{ and } \sigma \models \varphi_2, \\
\sigma &\models \varphi_1 \vee \varphi_2 \text{ iff } \sigma \models \varphi_1 \text{ or } \sigma \models \varphi_2, \\
\sigma &\models \neg \varphi \text{ iff } \sigma \not\models \varphi, \\
\sigma &\models \bigcirc \varphi \text{ iff } \sigma[1] \models \varphi, \\
\sigma &\models \Box \varphi \text{ iff } \forall 0 \leq i \leq n(\sigma[i] \models \varphi), \\
\sigma &\models \Diamond \varphi \text{ iff } \exists 0 \leq i \leq n(\sigma[i] \models \varphi), \\
\sigma &\models \varphi_1 \mathcal{A} \varphi_2 \text{ iff } \forall 0 \leq i < n(\text{if } \sigma[i] \models \varphi_1, \sigma[i+1] \models \varphi_2).
\end{aligned}$$

For example, we are given an LTL constraint ψ and three trajectories (σ_1 , σ_2 and σ_3) as shown in Figure 3. Clearly, we have $\sigma_1 \models \psi$ because, after public transit in S_2 and S_3 , traveling by car has never been taken place. Moreover, we have $\sigma_2 \not\models \psi$ because we have $M = c$ hold in S_6 and S_7 after having $M = p$ hold in S_2 and S_5 . Finally, we know $\sigma_3 \models \psi$, as the mode is always neither biking nor public transit.




Figure 3: State transition diagram

Preferences

A state is described as a set of *state variables*. The state variables of a state S include the transportation mode M that led to S , time T spent so far per mode (e.g., T_{bike} for biking and T_{public} for public transit), fare D spent so far per mode (e.g., D_{gas} for driving and D_{taxi} for taking a cab), and variables related to the auxiliary data once uploaded. These extra data related variables are metrics such as the sum (A_{sum}), the maximum (A_{max}), the minimum (A_{min}), and the average (A_{avg}) data along the path.

We focused on weighted functions over state variables and designed the cost function, called preferential cost function (PCF), that guides the graph-based search engine in our trip planner as follows.

$$\begin{aligned}
PCF(S) = & \beta_1 * (\alpha_{walk} \cdot T_{walk} + \alpha_{bike} \cdot T_{bike} + \dots) \\
& + (D_{gas} + D_{public} + \dots) \\
& + \beta_2 * (A_{sum} + \dots),
\end{aligned} \tag{3}$$

where α_i are real numbers representing the relations among different time pieces, and β_1 (β_2) is the ratio that essentially describes how much in dollars a user would pay to save an hour (an auxiliary data, respectively).

Preference Elicitation To gather these coefficients (α_i 's and β_i 's) in our *PCF*, we designed interface to elicit these numbers from the user. It asks the user questions and collect answers from the user to derive the coefficients. These questions are as follows.

1. One hr of walking = _____ hrs of driving?
2. One hr of biking = _____ hrs of driving?
3. One hr of public transit = _____ hrs of driving?
4. One hr of taxi = _____ hrs of driving?
5. How much in dollars would you pay to save an hour in traveling?
6. How much in dollars would you pay to avoid an auxiliary event (e.g., crime) in traveling?

For instance, Alice, an agent, answers 3, 2, 0.25 and 0.5 to the first four questions in the list above. Intuitively, the numbers indicate that she prefers public transit the most, followed by taxi, driving, biking and walking, in order. We show how we can now derive α_i 's in Equation 3. We start with setting $\alpha_{car} = 1$. Now, since one hour of walking is equivalent to 3 hours of walking, we have $\alpha_{walk} \times 1 = \alpha_{car} \times 3$; hence, we derive $\alpha_{walk} = 3$. Similarly, we have $\alpha_{bike} = 2$, $\alpha_{public} = 0.25$, and $\alpha_{taxi} = 0.5$. Computing β_1 and β_2 is straightforward. Indeed, function *PCF* with the input of time, fare and auxiliary metric pieces boils down to monetary cost, and the planner computes the best path by optimize based on this overall monetary cost in the searching process.

Reasoning with Constraints and Preferences




Figure 4: Adjusted A*

We leveraged the widely-used A* search algorithm on top of our high-performance graph search engine (cf. Figure 4). The A* algorithm incorporates the following cost function.

$$f(S) = g(S) + h(S), \tag{4}$$

where $g(S)$ is the cost of an optimal trip from the initial state to S , and $h(S)$ is an admissible estimate of the cost of an optimal trip from S to goal.

To prune the search space, we check satisfiability of the temporal constraints in LTL at expansion of the search tree. To guide the search engine, we set $g(S) = PCF(S)$ and $h(S)$ the minimum estimate among all available modes in S .

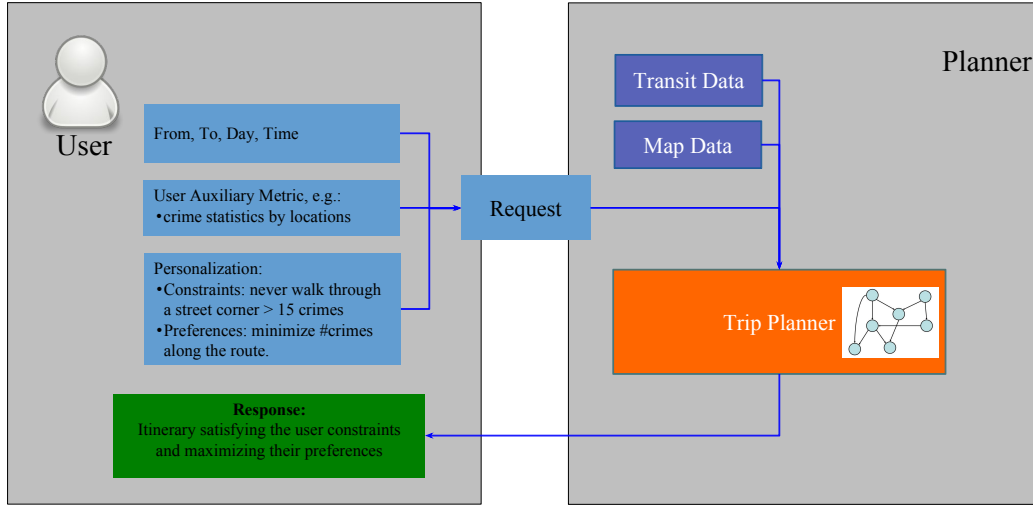


Figure 5: System Overview

Implementation and Results

System Architecture

We designed and implemented a multi-modal trip planning system (cf. Figure 5) based on a high-performance graph search engine. The planner allows user uploads, as well as declarative constraints and preferences. We now describe the structure of the planning system.

The trip planner takes two types of data as input: static data and user-specified request. The static input includes Map Data and Transit Data. Map Data describes the map, a directed graph where nodes are street corners, bus stops and train stations. Transit Data is a set of schedules for the buses and trains. On the other hand, a user provides her request, composed of three parts. First, the user enters *from* and *to* locations on the map together with day and time of the start of the trip. Second, the user may upload her auxiliary metric dataset, e.g., crime rates. Lastly, the user specifies her constraints in LTL and preferences as a *PCF*. For example, the constraint could be “never walk through a bad neighborhood.” Given these inputs, our planner computes an optimal path satisfying all the constraints and optimizing the preferences. Note that the request from the user is encapsulated into a JSON object.

For instance, the JSON object for the constraint ψ in Equation 2 is shown in Figure 6.

Results

We present and analyze resulting routes for three agents: Alice, Bob, and Cal. This is assuming no auxiliary metric datasets uploaded so that the agents focus on time and fare. Fixing their where and when information, we show how agents’ different constraints and preferences affect their optimal routes computed by our planner. Their where and when are set so that they all plan to travel from San Jose International Airport (SJC) in San Jose, to Pier 39 in San Francisco.

```
{
  "constraints": {
    "type": "after",
    "value": [
      {
        "type": "or",
        "value": [
          {
            "type": "=",
            "variable": "M",
            "value": "b"
          },
          {
            "type": "=",
            "variable": "M",
            "value": "p"
          }
        ]
      },
      {
        "type": "always",
        "value": {
          "type": "not",
          "value": {
            "type": "=",
            "variable": "M",
            "value": "c"
          }
        }
      }
    ]
  }
}
```

Figure 6: JSON object for constraint ψ in Equation 2

Note that the natural constraint ψ in Equation 2 is implicitly imposed on all cases.

Agent Alice has only one constraint that she does not have a car; thus, driving to her is never available during this trip. Per her preferences, Alice provides her thoughts as earlier; that is, public transportation is preferred to taxi, which is better than biking, preferred to walking. What’s more, she decides that she would sacrifice thirty dollars to save one hour during the trip.

The resulting route for Alice is included in Figure 7. The solution takes 2 hours 7 minutes and 18 dollars 94 cents. It is optimal in the sense the combined value of time and fare is minimal among all paths from SJC to Pier 39.

Like Alice, Bob is constrained that he will not drive a car in his travel, and his dollar per hour is thirty. Moreover, Bob makes his mind to have some workout with his bike and dic-

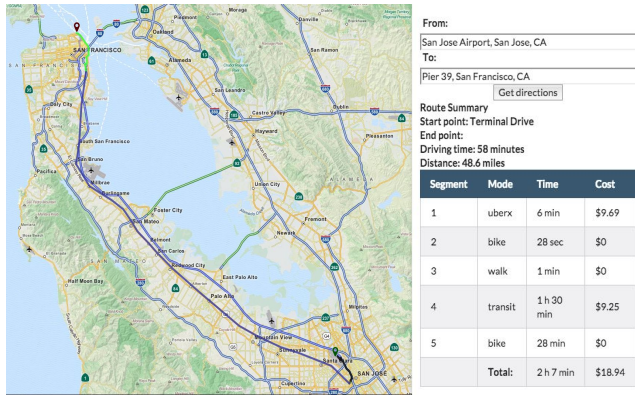


Figure 7: Resulting route for Alice

tates that he will bike for between one and two hours. He expresses his preferences: biking and public transit are the most preferred, next is taxi, and the least preferred is walking. Similarly, he has done so by answering the aforementioned elicitation questions.

The result for Bob is depicted in Figure 8. It spans 2 hours 57 minutes in time with the fare of 29 dollars 17 cents. It is so, seemingly worse than what Alice achieved, only because Bob has the constraint that he will for sure bike for one to two hours, and the preferences that put biking the most satisfying mode.

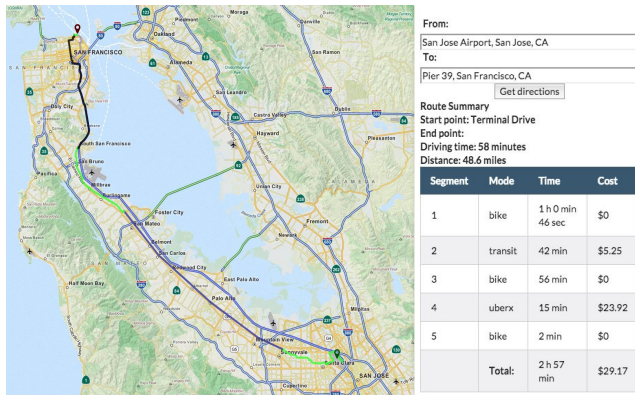


Figure 8: Resulting route for Bob

Agent Cal will also not use a car in the trip. He affirms that time is greatly valuable and one hour is measured to 500 dollars, although his budget is restricted to 50. Cal's preferences are that the most preferred are public transit and taxi, and the next preferred are walking and biking that are equivalent.

Refer to Figure 9 for the optimal path for Cal. This route, stretching 1 hour 48 minutes, is the most time-saving one compared with the previous two, at the price of 49 dollars 91 cents. This is due to the constraint that Cal can spend up to 50 dollars, as well as his preferences and value of time being high.

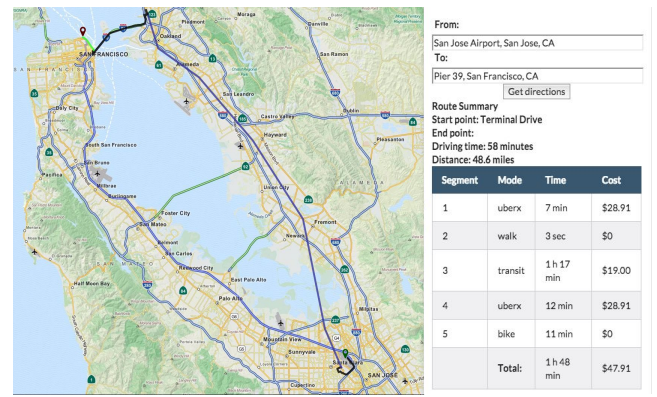


Figure 9: Resulting route for Cal

Auxiliary Metric When the user of the planner is interested in metrics other than the ones offered already (i.e., time and fare), she might discover new metrics (e.g., crime rates and pollution statistics), upload them into the planner, and retrieve optimal plans taking these metrics into account. One scenario of this approach is the following.

Our agent needs to travel without a car across San Francisco downtown at night. For her, safety is important. Having found the crime statistics for the area, the agent uploads the data as a new auxiliary metric into the map. By specifying that she will never walk through a neighborhood with more than fifteen crimes over the last month, and that she would sacrifice a quarter to avoid one crime incident, the agent tells the planner to come up with a relatively safe route. An example is shown in Figure 10, where the agent needs to start at the east of downtown and travel across the area to arrive at the west side. The computed path is represented by the line colored by black, blue and green, denoting taxi, public transit and biking, respectively. Clearly, this path routes away from crime-heavy areas and achieves optimality in that the combined metrics – time, fare and crime rates, uploaded and personalized by the user – is minimal among all possibilities.




Figure 10: Optimal route considering crime rates

Conclusion and Future Work

In multi-modal trip planning, it is vital for the planner to enable users to integrate new metrics of interest into the existing map, and to express constraints and preferences over which optimal routes are computed. In our work, we designed and built a graph-search based, extensible and personalizable multi-modal trip planner that utilizes linear temporal logic and preferential cost function to model user constraints and preferences.

As a note on the future work, we plan to explore techniques in planning for computing multiple routes that are diverse with bounded difference of costs. Continuing this direction, we intend to study the problem of learning the *PCF* coefficients using the observations of the decisions the user made among the computed paths. (It is reasonable to assume that routes picked by the user are preferred to others, and these observations could contribute to the learning problem.) Also interesting is to introduce traffic information into the planner to support real-time multi-agent concurrent trip planning.

References

- [Antoniou, Matsoukis, and Roussi 2007] Antoniou, C.; Matsoukis, E.; and Roussi, P. 2007. A methodology for the estimation of value-of-time using state-of-the-art econometric models. *Journal of Public Transportation* 10(3):1.
- [Bast et al. 2015] Bast, H.; Delling, D.; Goldberg, A.; Müller-Hannemann, M.; Pajor, T.; Sanders, P.; Wagner, D.; and Werneck, R. F. 2015. Route planning in transportation networks. *arXiv preprint arXiv:1504.05140*.
- [Bienvenu, Fritz, and McIlraith 2011] Bienvenu, M.; Fritz, C.; and McIlraith, S. A. 2011. Specifying and computing preferred plans. *Artificial Intelligence* 175(7):1308–1345.
- [Modesti and Sciomachen 1998] Modesti, P., and Sciomachen, A. 1998. A utility measure for finding multiobjective shortest paths in urban multimodal transportation networks. *European Journal of Operational Research* 111(3):495–508.
- [Son and Pontelli 2004] Son, T. C., and Pontelli, E. 2004. Planning with preferences using logic programming. In *Logic Programming and Nonmonotonic Reasoning*. Springer. 247–260.
- [Zhou and Hansen 2011] Zhou, R., and Hansen, E. A. 2011. Dynamic state-space partitioning in external-memory graph search. In *ICAPS*.